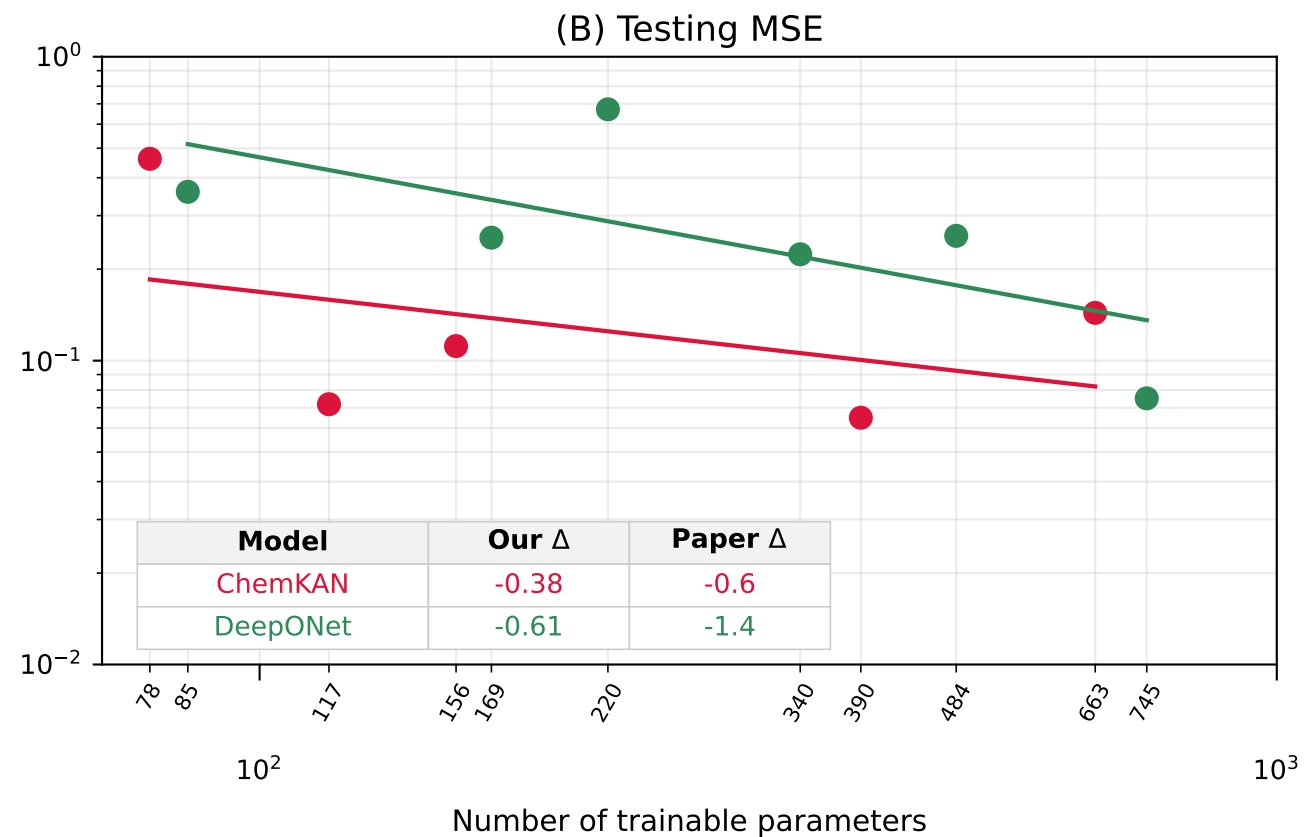
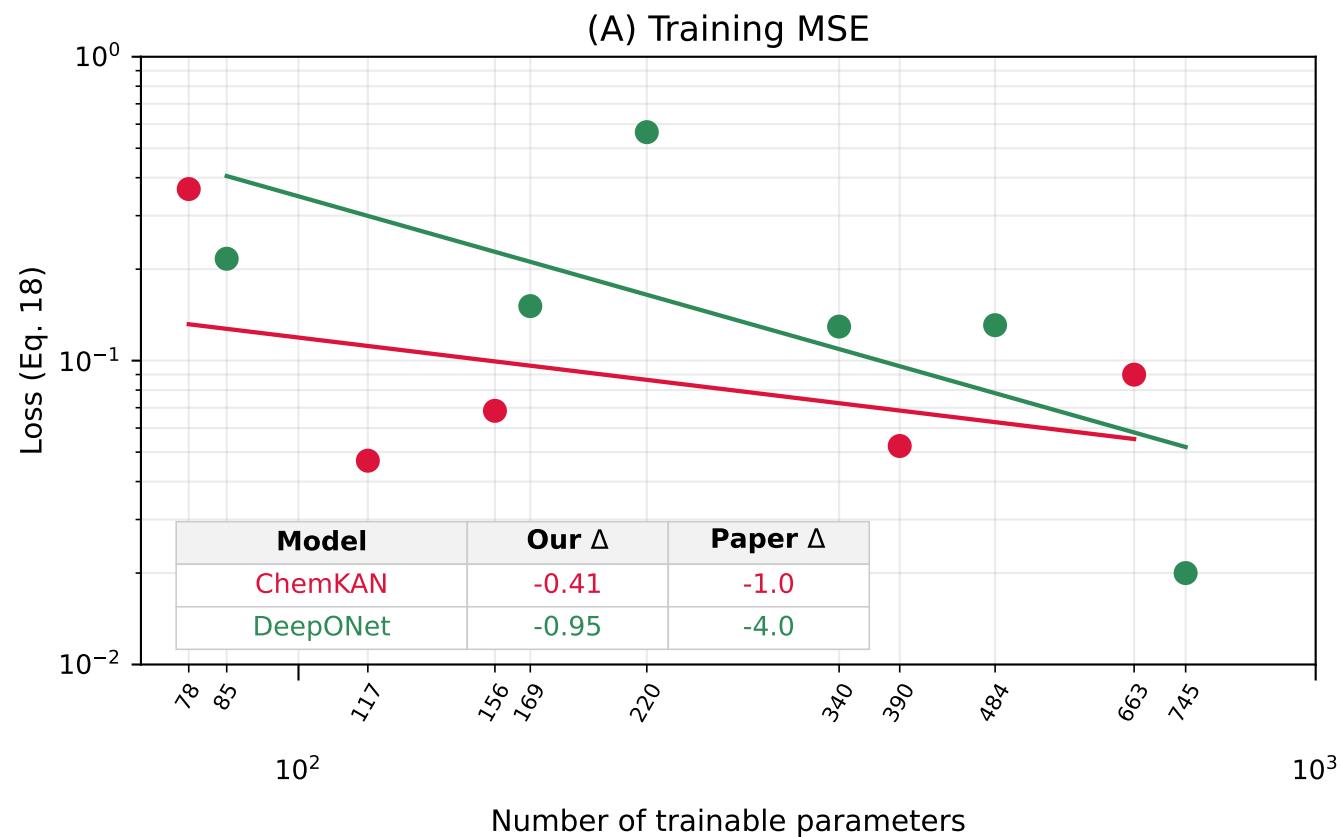


Fig. 4 - neural scaling: our results and paper slope comparison
 Circles and solid lines: our results and fitted trends. The tables compare our slopes with the paper's reported slopes.



Δ is the slope of $\log(\text{loss})$ vs $\log(\text{parameters})$. Our fits use all runs; the paper reports fits before saturation.